

```
In [1]: import sagemaker
import boto3
from sagemaker.amazon.amazon_estimator import get_image_uri
from sagemaker.session import s3_input, Session
```

/home/ec2-user/anaconda3/envs/python3/lib/python3.10/site-packages/pandas/core/computation/expressions.py:21: UserWarning: Pandas requires version '2.8.0' or newer of 'numexpr' (version '2.7.3' currently installed).
from pandas.core.computation.check import NUMEXPR_INSTALLED

sagemaker.config INFO - Not applying SDK defaults from location: /etc/xdg/sagemaker/config.yaml
sagemaker.config INFO - Not applying SDK defaults from location: /home/ec2-user/.config/sagemaker/config.yaml

```
In [2]: bucketname = 'praveen-sage-maker-practice-2'
my_region = boto3.session.Session().region_name
print(my_region)
```


us-east-1

```
In [3]: s3 = boto3.resource('s3')
s3.create_bucket(Bucket = bucketname)
```

```
Out[3]: s3.Bucket(name='praveen-sage-maker-practice-2')
```

```
In [4]: prefix = 'xgboost-as-abuilt-in-algo'
output_path = 's3://{}/{}/output'.format(bucketname, prefix)
```

```
In [5]: import pandas as pd
import urllib
urllib.request.urlretrieve("https://d1.awsstatic.com/tmt/build-train-deploy-ma
model_data = pd.read_csv('./bank_clean.csv', index_col = 0)
```




In [6]: `model_data.head()`

Out[6]:

	age	campaign	pdays	previous	no_previous_contact	not_working	job_admin.	job_blue-collar	j
0	56	1	999	0	1	0	0	0	
1	57	1	999	0	1	0	0	0	
2	37	1	999	0	1	0	0	0	
3	40	1	999	0	1	0	1	0	
4	56	1	999	0	1	0	0	0	

5 rows × 61 columns




In [7]: `import numpy as np`
`train_data, test_data = np.split(model_data.sample(frac=1, random_state=1729),`
`print(train_data.shape, test_data.shape)`

(28831, 61) (12357, 61)

/home/ec2-user/anaconda3/envs/python3/lib/python3.10/site-packages/numpy/core/fromnumeric.py:57: FutureWarning: 'DataFrame.swapaxes' is deprecated and will be removed in a future version. Please use 'DataFrame.transpose' instead.

return bound(*args, **kws)

In [8]: `import os`
`pd.concat([train_data['y_yes'], train_data.drop(['y_yes', 'y_no'], axis = 1)], axis=1)`



In [9]: `boto3.Session().resource('s3').Bucket(bucketname).Object(os.path.join(prefix,`
`s3_input_train = sagemaker.TrainingInput(s3_data='s3://{}/{}/train'.format(bucketname,`

In [10]: `# Test Data Into Buckets`
`pd.concat([test_data['y_yes'], test_data.drop(['y_no', 'y_yes'], axis=1)], axis=1)`
`boto3.Session().resource('s3').Bucket(bucketname).Object(os.path.join(prefix,`
`s3_input_test = sagemaker.TrainingInput(s3_data='s3://{}/{}/test'.format(bucketname,`

In [11]: `container = get_image_uri(boto3.Session().region_name, 'xgboost', repo_version=`

The method `get_image_uri` has been renamed in `sagemaker>=2`.

See: <https://sagemaker.readthedocs.io/en/stable/v2.html> (<https://sagemaker.readthedocs.io/en/stable/v2.html>) for details.

```
In [12]: hyperparameters = {
        "max_depth": "5",
        "eta": "0.2",
        "gamma": "4",
        "min_child_weight": "6",
        "subsample": "0.7",
        "objective": "binary:logistic",
        "num_round": 50
    }
```

```
In [14]: estimator = sagemaker.estimator.Estimator(image_uri=container,
        hyperparameters=hyperparameters,
        role=sagemaker.get_execution_role(),
        train_instance_count=1,
        train_instance_type='ml.m5.2xlarge',
        train_volume_size=5, # 5 GB
        output_path=output_path,
        train_use_spot_instances=True,
        train_max_run=300,
        train_max_wait=600)
```

sagemaker.config INFO - Not applying SDK defaults from location: /etc/xdg/sagemaker/config.yaml

sagemaker.config INFO - Not applying SDK defaults from location: /home/ec2-user/.config/sagemaker/config.yaml

train_instance_count has been renamed in sagemaker>=2.

See: <https://sagemaker.readthedocs.io/en/stable/v2.html> (<https://sagemaker.readthedocs.io/en/stable/v2.html>) for details.

train_instance_type has been renamed in sagemaker>=2.

See: <https://sagemaker.readthedocs.io/en/stable/v2.html> (<https://sagemaker.readthedocs.io/en/stable/v2.html>) for details.

train_max_run has been renamed in sagemaker>=2.

See: <https://sagemaker.readthedocs.io/en/stable/v2.html> (<https://sagemaker.readthedocs.io/en/stable/v2.html>) for details.

train_use_spot_instances has been renamed in sagemaker>=2.

See: <https://sagemaker.readthedocs.io/en/stable/v2.html> (<https://sagemaker.readthedocs.io/en/stable/v2.html>) for details.

train_max_wait has been renamed in sagemaker>=2.

See: <https://sagemaker.readthedocs.io/en/stable/v2.html> (<https://sagemaker.readthedocs.io/en/stable/v2.html>) for details.

train_volume_size has been renamed in sagemaker>=2.

See: <https://sagemaker.readthedocs.io/en/stable/v2.html> (<https://sagemaker.readthedocs.io/en/stable/v2.html>) for details.

sagemaker.config INFO - Not applying SDK defaults from location: /etc/xdg/sagemaker/config.yaml

sagemaker.config INFO - Not applying SDK defaults from location: /home/ec2-user/.config/sagemaker/config.yaml

```
In [15]: estimator.fit({'train':s3_input_train,'validation':s3_input_test})
```

```
INFO:sagemaker:Creating training-job with name: sagemaker-xgboost-2023-12-12-00-06-01-158
```

```
2023-12-12 00:06:01 Starting - Starting the training job...
```

```
2023-12-12 00:06:14 Starting - Insufficient capacity error from EC2 while launching instances, retrying
```

```
g!.....
```

```
2023-12-12 00:16:15 Stopping - Stopping the training job
```

```
2023-12-12 00:16:15 MaxWaitTimeExceeded - Training job wait time exceeded MaxWaitTimeInSeconds provided
```

```
..
```

```
WARNING:sagemaker:Job ended with status 'Stopped' rather than 'Completed'. This could mean the job timed out or stopped early for some other reason: Consider checking whether it completed as you expect.
```

```
In [*]: sagemaker.Session().delete_endpoint(xgb_predictor.endpoint)
        bucket_to_delete = boto3.resource('s3').Bucket(bucket_name)
        bucket_to_delete.objects.all().delete()
```

```
In [ ]:
```